

Winning Space Race with Data Science

Taeef Najib (<https://github.com/taeefnajib>)
April 10, 2022



Outline

- Executive Summary - 3
- Introduction - 4
- Methodology - 5
- Results
- Conclusion
- Appendix

Executive Summary

- We have collected data from public SpaceX API and by scrapping SpaceX Wikipedia page. Created labels column 'Class' which denotes all the successful landings. We performed EDA (Exploratory Data Analysis) using SQL, visualization, folium maps, and dashboards. Then we gathered relevant columns to be used as our features. After that we performed One Hot Encoding and changed all the categorical variables to binary. We standardized data and used GridSearchCV to find the best parameters for our machine learning models. Finally, we visualized the accuracy score of all models.
- We used four machine learning models- Logistic Regression, Support Vector Machine (SVM), Decision Tree Classifier, and K Nearest Neighbors. All produced similar results with the accuracy rate of almost 83.33%. All models over predicted successful landings. It is clear that more data is required for a better model determination and accuracy.

Introduction

Background and Context

- In this capstone, we will predict if the Falcon 9 first stage will land successfully. SpaceX advertises Falcon 9 rocket launches on its website, with a cost of 62 million dollars; other providers cost upward of 165 million dollars each, much of the savings is because SpaceX can reuse the first stage. Therefore if we can determine if the first stage will land, we can determine the cost of a launch. This information can be used if an alternate company wants to bid against SpaceX for a rocket launch. In this module, you will be provided with an overview of the problem and the tools you need to complete the course. Space Y wants to compete with Space X

Problem at Hand

- Space Y tasks us to train a machine learning model to predict successful Stage 1 recovery

Section 1

Methodology

Methodology

Executive Summary

- Data collection methodology:
 - Gathered data from SpaceX public API and by scrapping SpaceX Wikipedia page
- Perform data wrangling
 - Classifying true landings as successful and unsuccessful otherwise
- Perform exploratory data analysis (EDA) using visualization and SQL
- Perform interactive visual analytics using Folium and Plotly Dash
- Perform predictive analysis using classification models
 - We tuned the models using GridSearchCV

Data Collection

Data collection process involved a combination of API requests from Space X public API and web scraping data from a table in Space X's Wikipedia entry.

The next slide will show the flowchart of data collection from API and the one after will show the flowchart of data collection from webscraping.

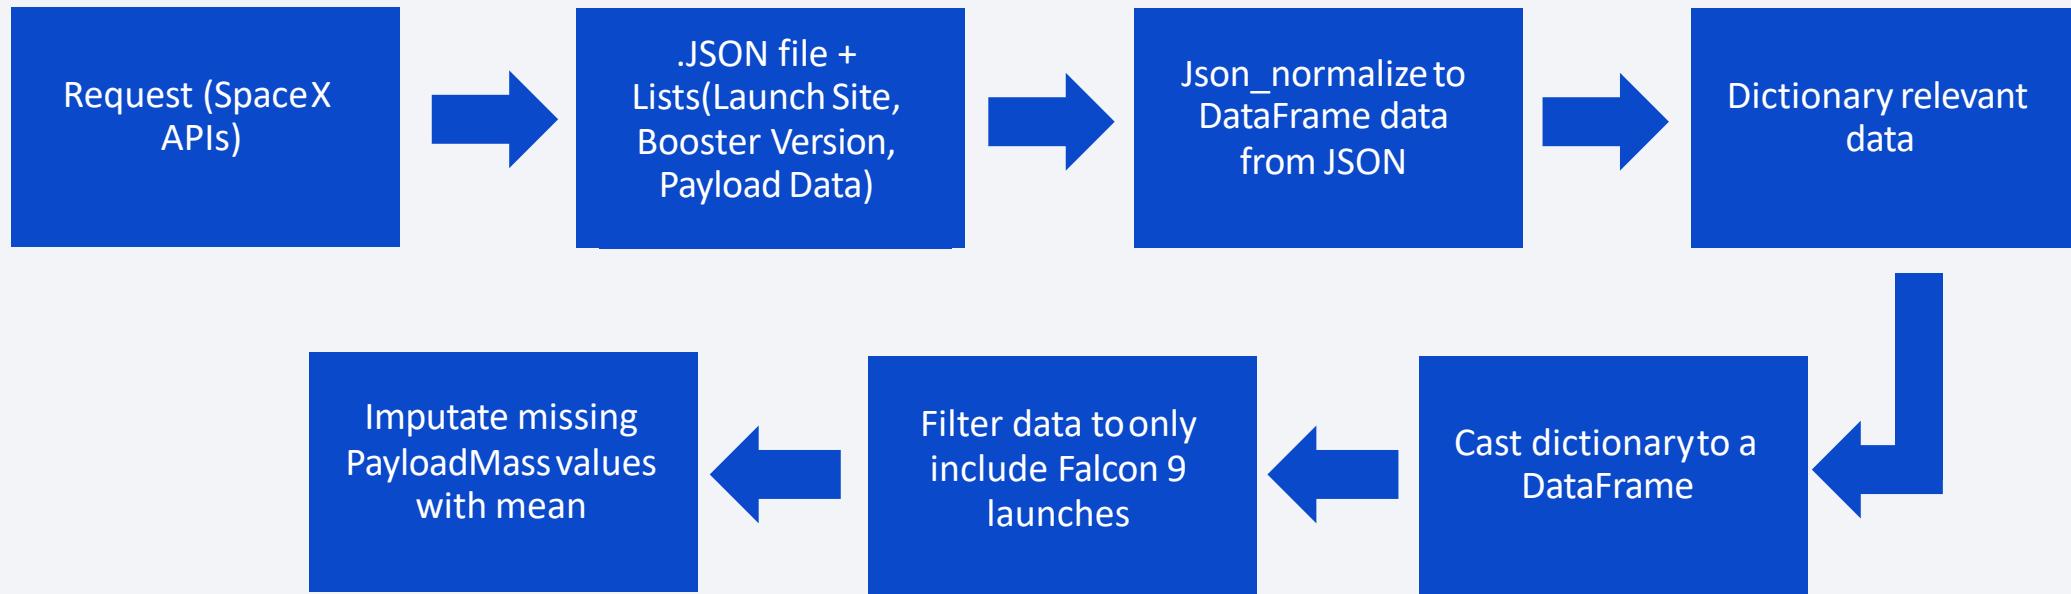
Space X API Data Columns:

FlightNumber, Date, BoosterVersion, PayloadMass, Orbit, LaunchSite, Outcome, Flights, GridFins, Reused, Legs, LandingPad, Block, ReusedCount, Serial, Longitude, Latitude

Wikipedia Webscrape Data Columns:

Flight No., Launch site, Payload, PayloadMass, Orbit, Customer, Launch outcome, Version Booster, Booster landing, Date, Time

Data Collection – SpaceX API



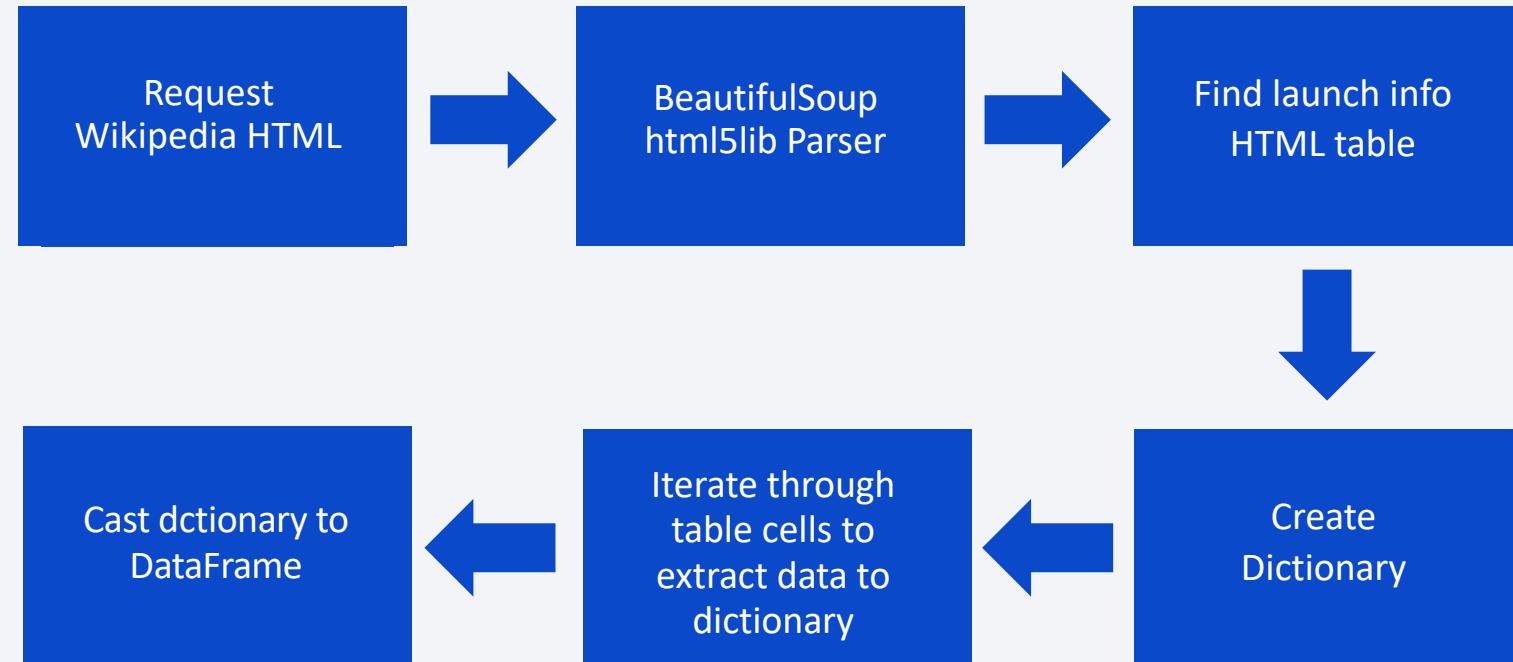
GitHub URL:

<https://github.com/taeefnajib/IBM-Applied-Data-Science-Capstone/blob/main/Data%20Collection%20Using%20API.ipynb>

Data Collection - Scraping

GitHub URL:

<https://github.com/taeefnajib/IBM-Applied-Data-Science-Capstone/blob/main/Data%20Collection%20Using%20API.ipynb>



Data Wrangling

Create a training label with landing outcomes where successful = 1 & failure = 0.

Outcome column has two components: 'Mission Outcome' 'Landing Location'

New training label column 'class' with a value of 1 if 'Mission Outcome' is True and 0 otherwise.

Value Mapping:

True ASDS, True RTLS, & True Ocean – set to -> 1

None None, False ASDS, None ASDS, False Ocean, False RTLS – set to -> 0

GitHub URL:

<https://github.com/taeefnajib/IBM-Applied-Data-Science-Capstone/blob/main/Data%20Wrangling.ipynb>

EDA with Data Visualization

Exploratory Data Analysis performed on variables Flight Number, Payload Mass, Launch Site, Orbit, Class and Year.

Plots Used:

Flight Number vs. Payload Mass, Flight Number vs. Launch Site, Payload Mass vs. Launch Site, Orbit vs. Success Rate, Flight Number vs. Orbit, Payload vs Orbit, and Success Yearly Trend

Scatter plots, line charts, and bar plots were used to compare relationships between variables to decide if a relationship exists so that they could be used in training the machine learning model

GitHub URL:

<https://github.com/taeefnajib/IBM-Applied-Data-Science-Capstone/blob/main/EDA%20with%20Visualization.ipynb>

EDA with SQL

- Loaded data set into IBM DB2 Database.
- Queried using SQL Python integration.
- Queries were made to get a better understanding of the dataset.
- Queried information about launch site names, mission outcomes, various pay load sizes of customers and booster versions, and landing outcomes

GitHub URL:

<https://github.com/taeefnajib/IBM-Applied-Data-Science-Capstone/blob/main/EDA%20with%20SQL.ipynb>

Build an Interactive Map with Folium

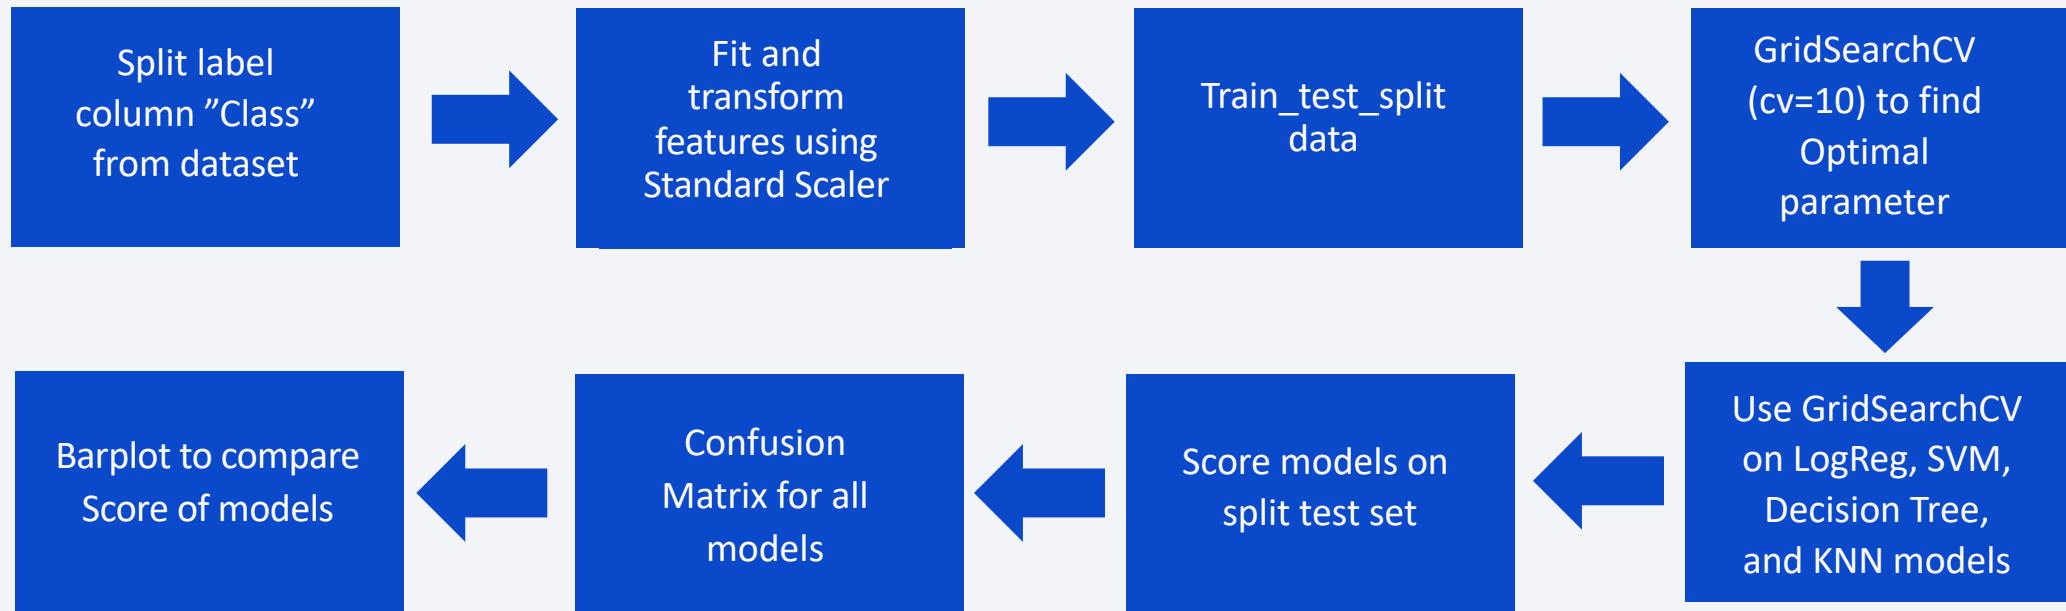
Folium maps mark Launch Sites, successful and unsuccessful landings, and a proximity example to key locations: Railway, Highway, Coast, and City.

This allows us to understand why launch sites may be located where they are. Also visualizes successful landings relative to location.

GitHub URL:

<https://github.com/taeefnajib/IBM-Applied-Data-Science-Capstone/blob/main/Interactive%20Visual%20Analytics%20and%20Dashboard.ipynb>

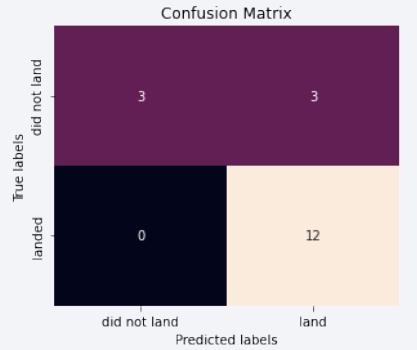
Predictive Analysis (Classification)



GitHub URL:

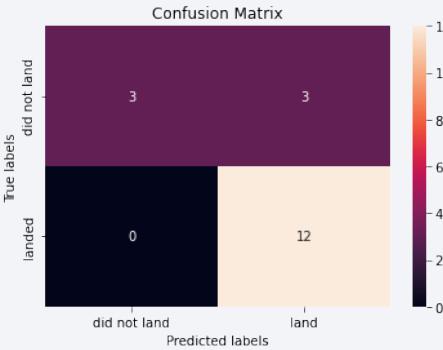
<https://github.com/taeefnajib/IBM-Applied-Data-Science-Capstone/blob/main/Predictive%20Analysis.ipynb>

Results



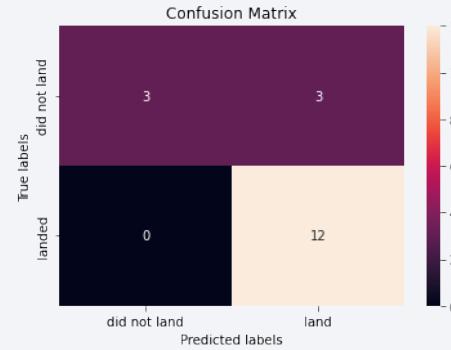
LogReg

Accuracy: 83.33%



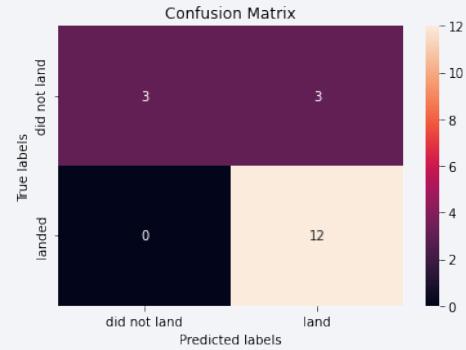
SVM

Accuracy: 83.33%



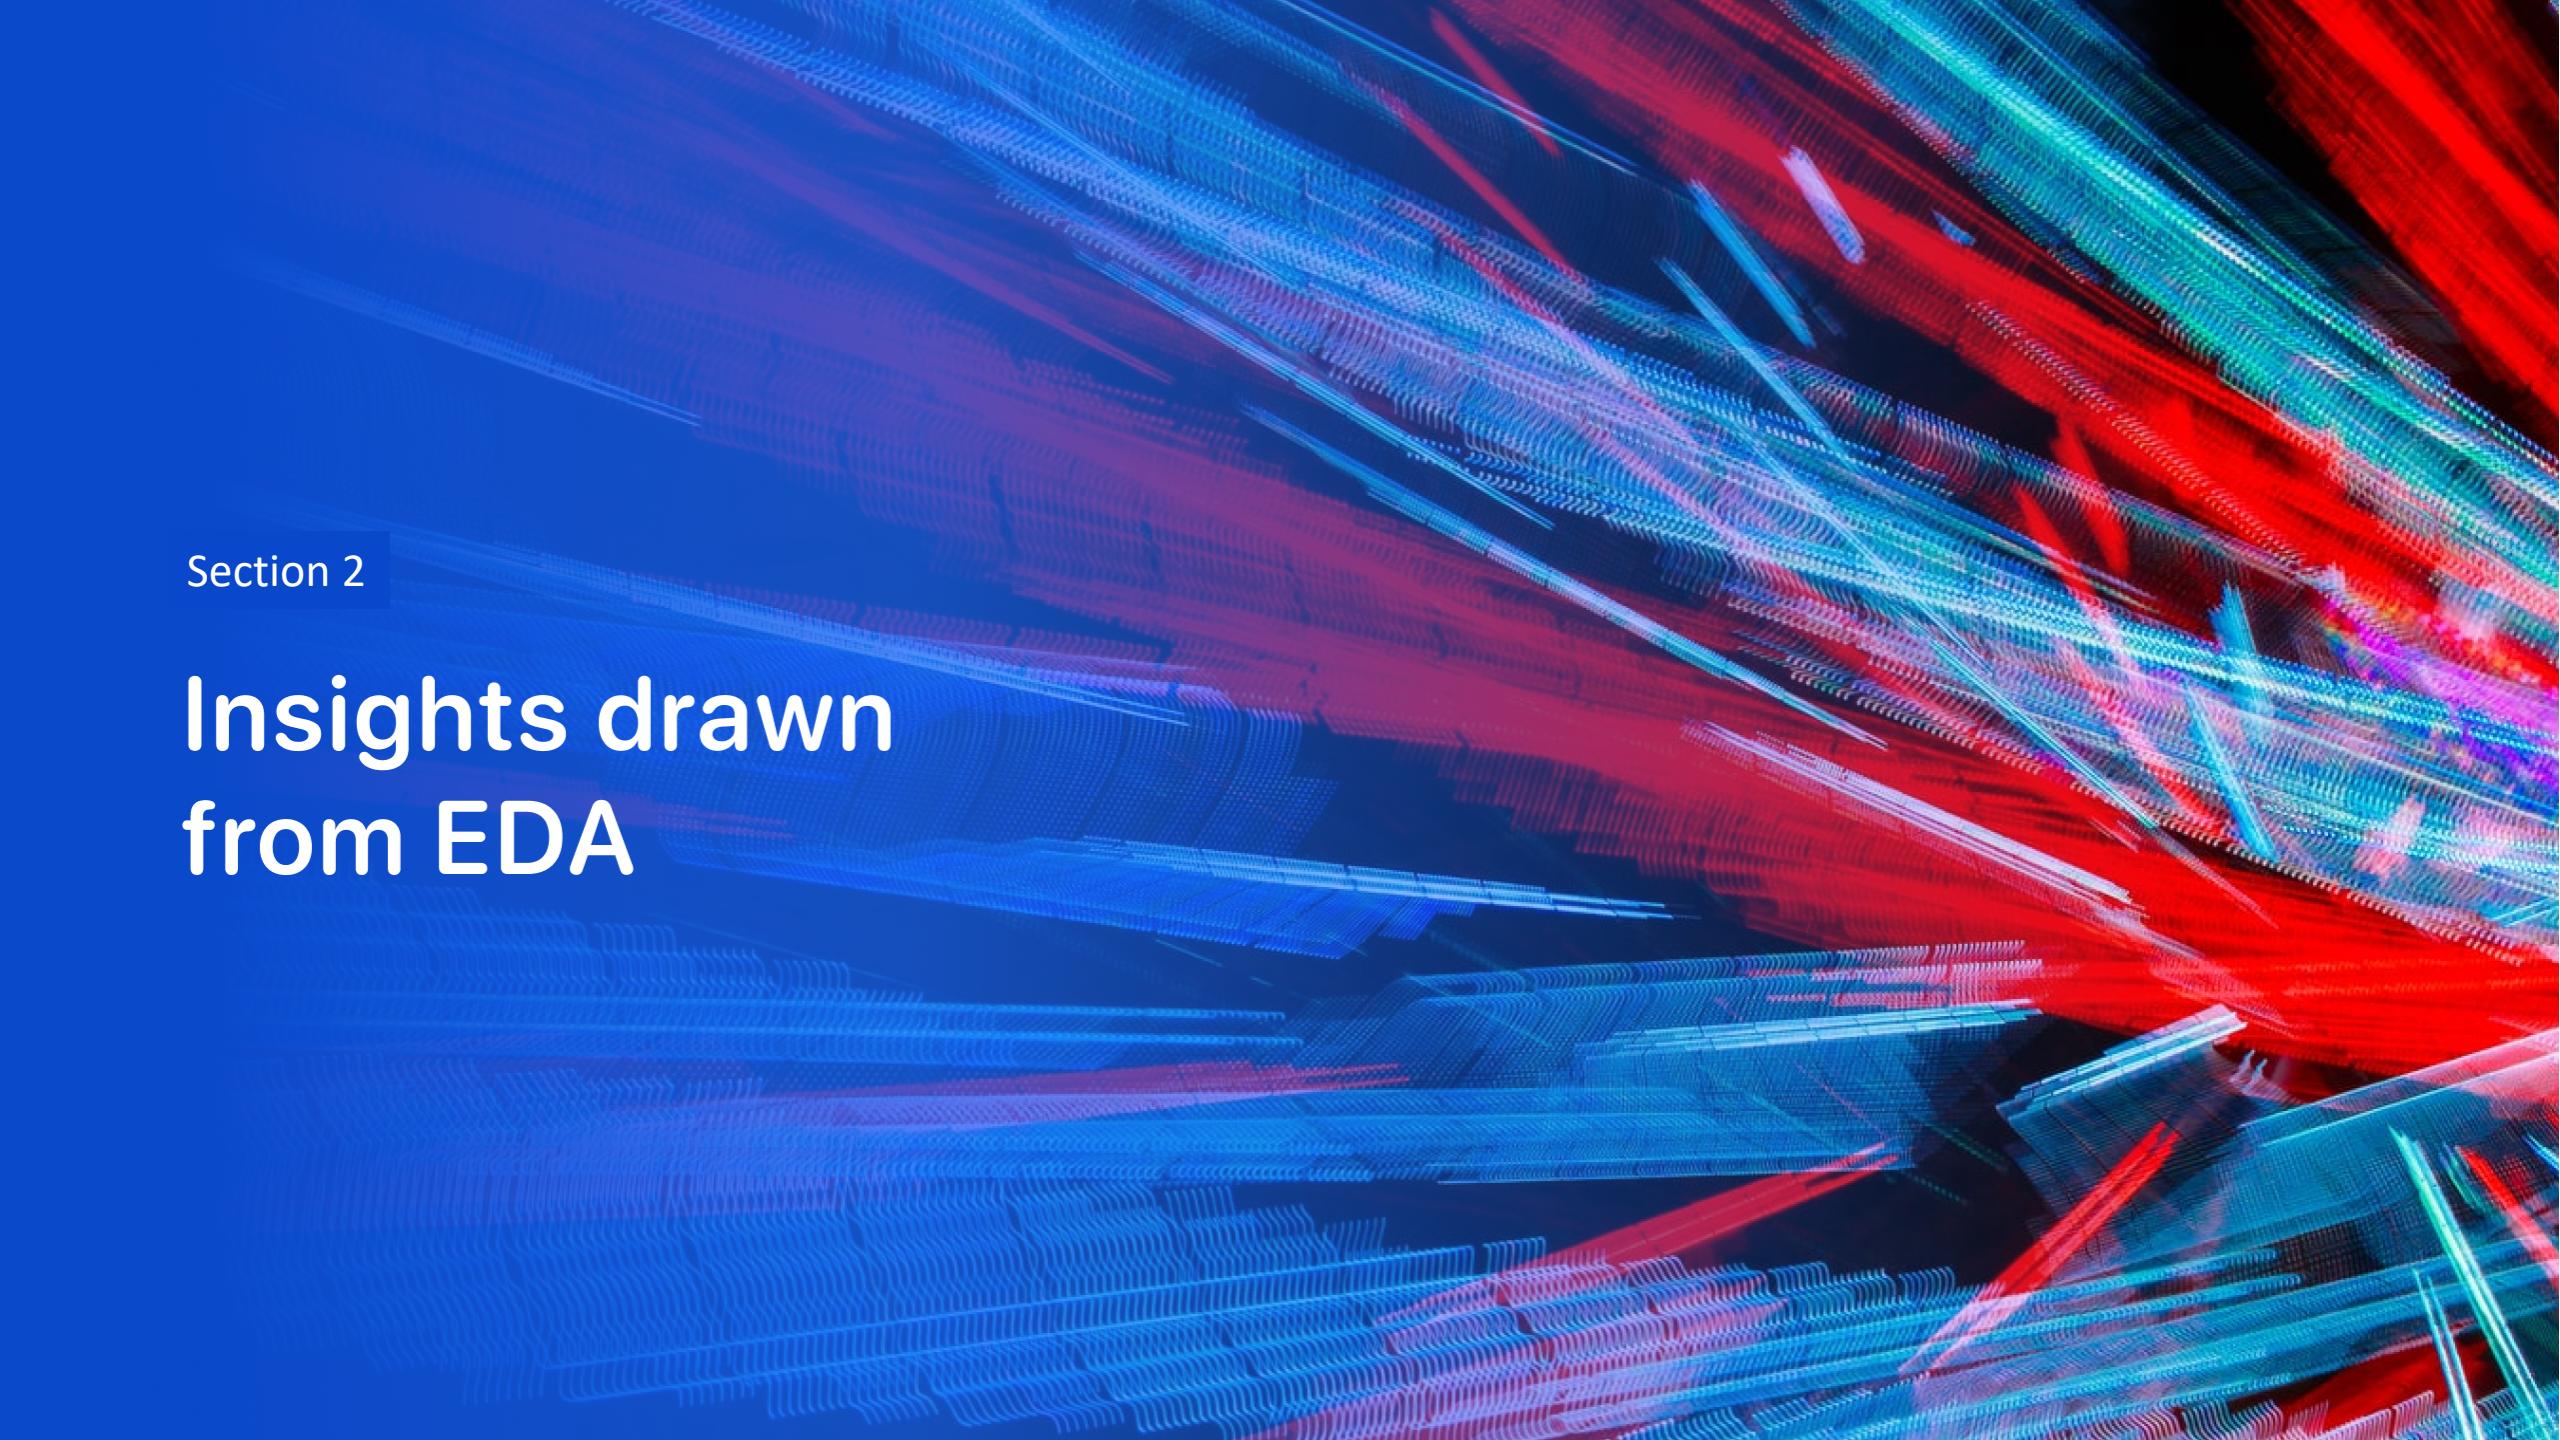
Decision Tree

Accuracy: 83.33%



KNN

Accuracy: 83.33%

The background of the slide features a complex, abstract pattern of glowing lines in shades of blue, red, and purple. These lines are arranged in a way that suggests depth and motion, resembling a 3D space filled with data or energy flow. The lines are thin and have a slight glow, creating a futuristic and high-tech feel.

Section 2

Insights drawn from EDA

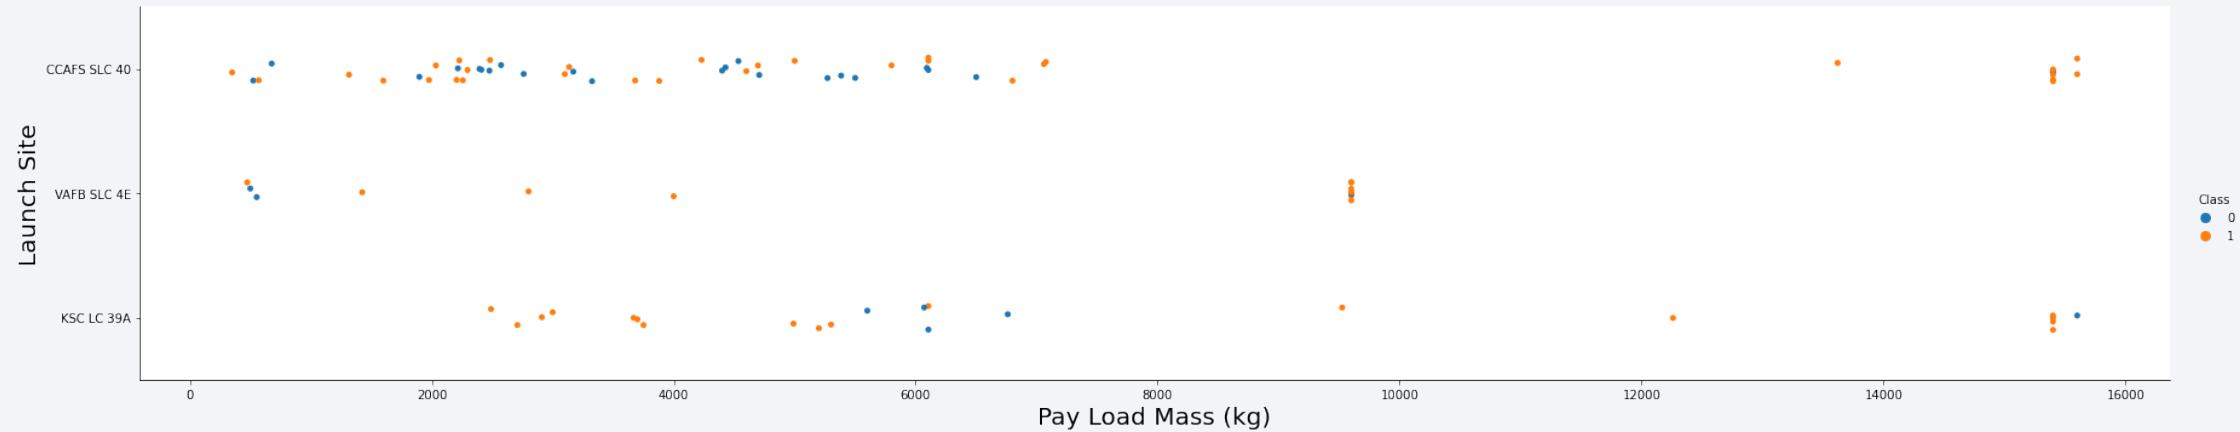
Flight Number vs. Launch Site



Blue indicates successful launch and orange indicates unsuccessful launch. Graphic suggests an increase in success rate over time (indicated in Flight Number). Likely a big breakthrough around flight 20 which significantly increased success rate.

CCAFS appears to be the main launch site as it has the most volume.

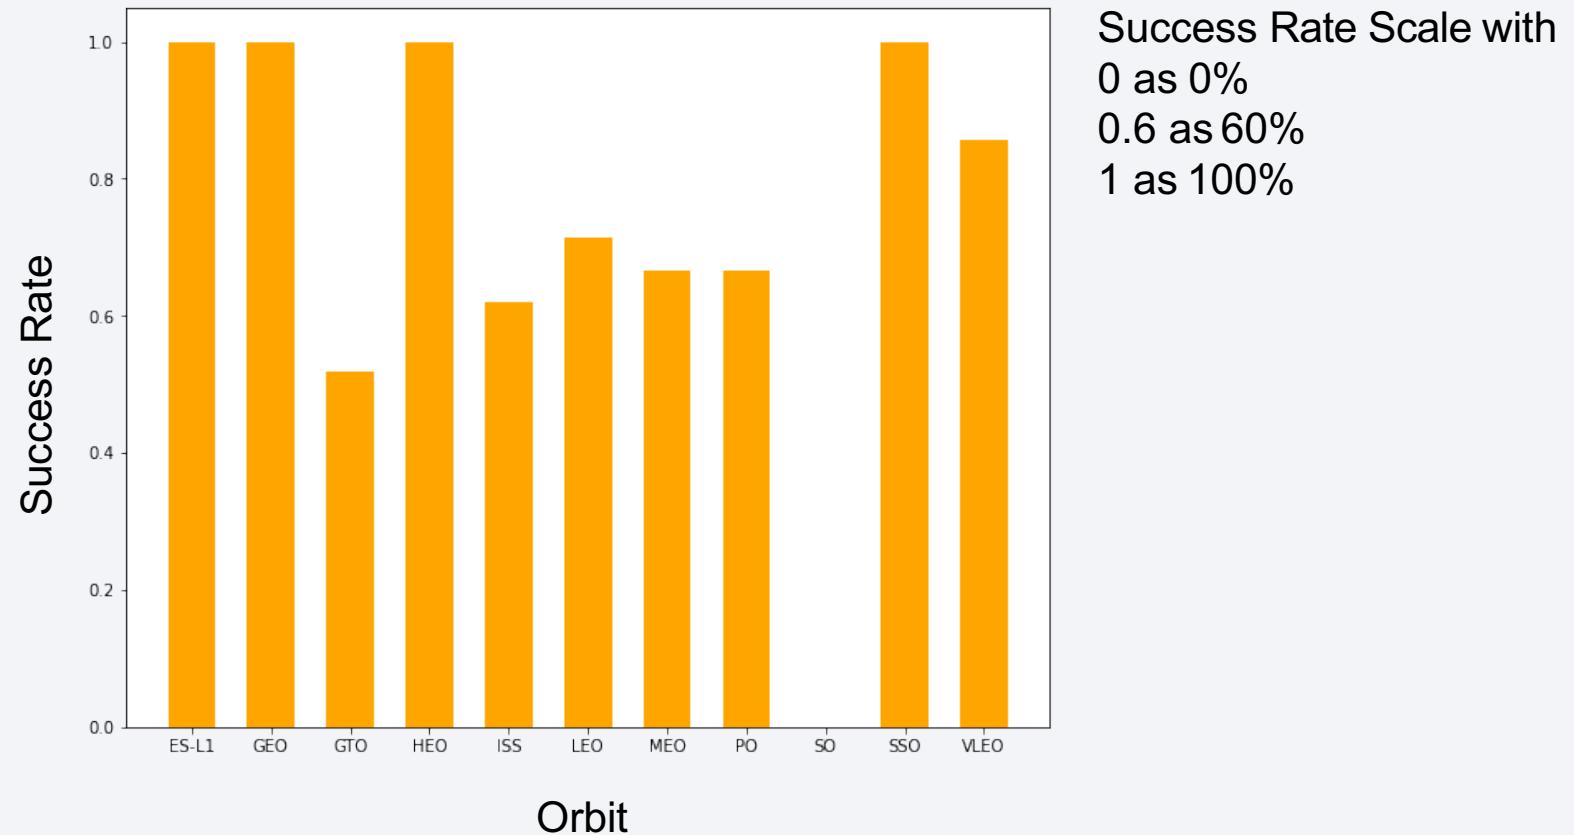
Payload vs. Launch Site



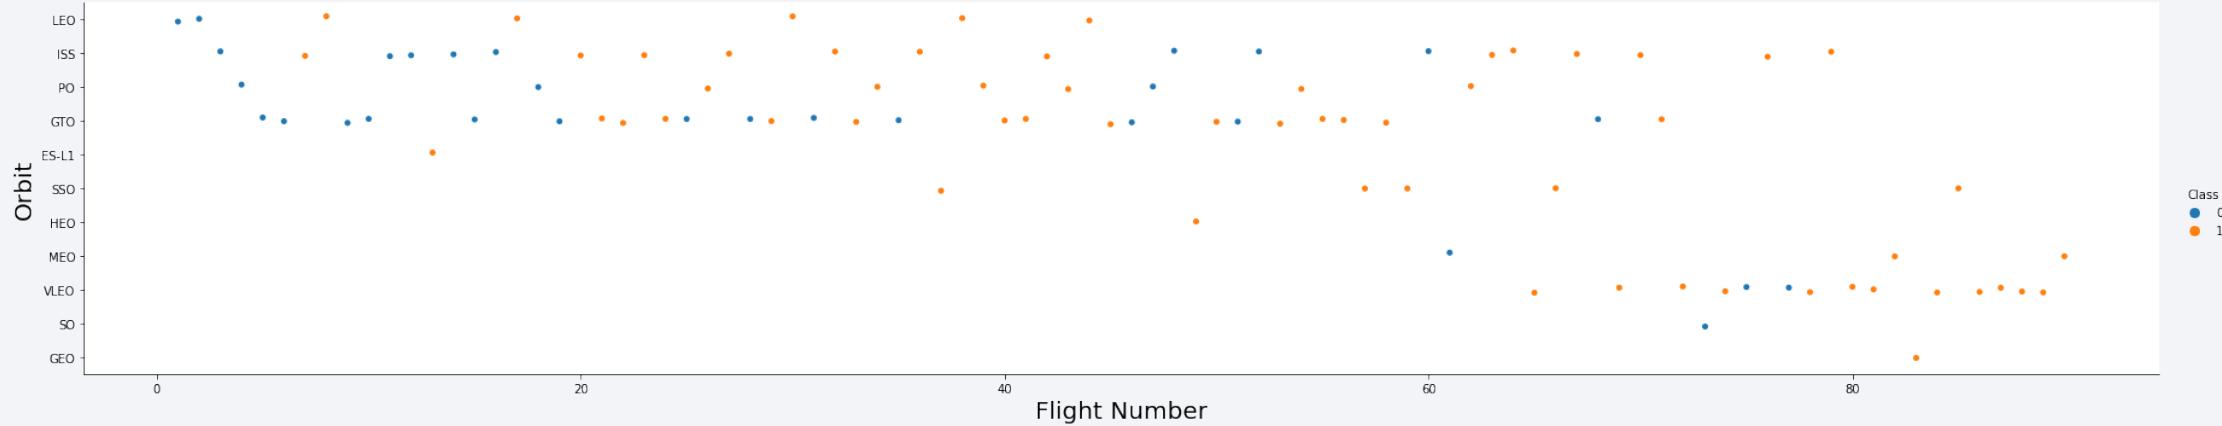
Blue indicates successful launch and orange indicates unsuccessful launch. Payload mass appears to fall mostly between 0-6000 kg. Different launch sites also seem to use different payload mass.

Success Rate vs. Orbit Type

ES-L1 (1), GEO (1), HEO (1)
have 100% success rate
(sample sizes in parenthesis)
SSO (5) has 100% success rate
VLEO (14) has decent success
rate and attempts
SO (1) has 0% success rate
GTO (27) has the around 50%
success rate but largest sample



Flight Number vs. Orbit Type



Blue indicates successful launch and orange indicates unsuccessful launch.

- Launch Orbit preferences changed over Flight Number.
- Launch Outcome seems to correlate with this preference.
- SpaceX started with LEO orbits which saw moderate success LEO and returned to VLEO in recent launches
- SpaceX appears to perform better in lower orbits or Sun-synchronous orbits

Payload vs. Orbit Type

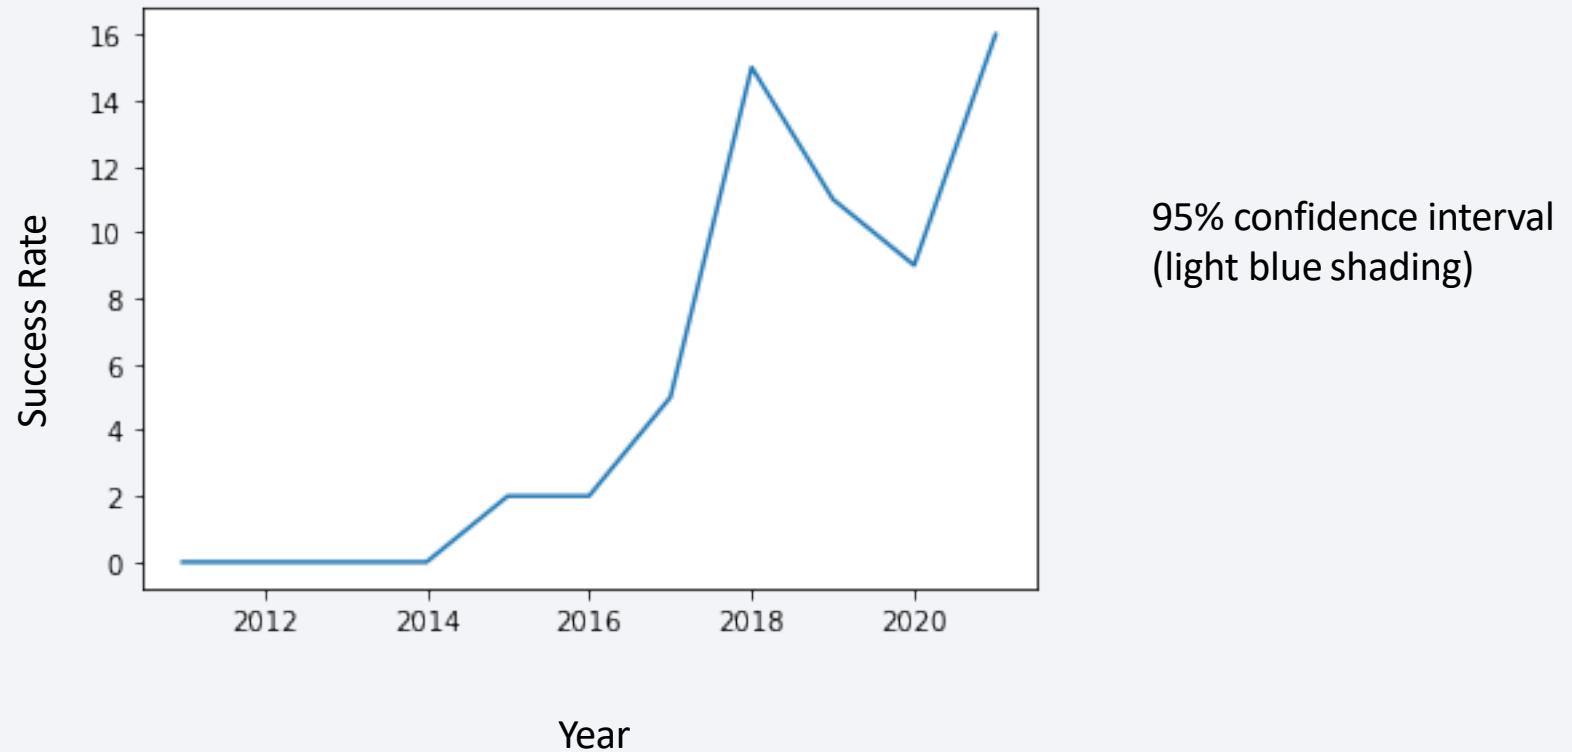


Blue indicates successful launch and orange indicates unsuccessful launch.

- Payload mass seems to correlate with orbit
- LEO and SSO seem to have relatively low payload mass
- The other most successful orbit VLEO only has payload mass values in the higher end of the range

Launch Success Yearly Trend

- Success generally increases over time since 2013 with a slight dip in 2018
- Success in recent years at around 80%



All Launch Site Names

- CCAFS LC-40
- CCAFS SCL-40
- KSC LC-39A
- VAFB SLC-4E

% %sql

```
SELECT DISTINCT
LAUNCH_SITE FROM
SPACEXDATASET;
```

launch_site

CCAFS LC-40

CCAFS SLC-40

KSC LC-39A

VAFB SLC-4E

Launch Site Names Begin with 'CCA'

% %sql

```
SELECT * FROM SPACEXDATASET WHERE LAUNCH_SITE LIKE 'CCA%' LIMIT 5;
```

DATE	time_utc_	booster_version	launch_site	payload	payload_mass_kg_	orbit	customer	mission_outcome	landing_outcome
2010-06-04	18:45:00	F9 v1.0 B0003	CCAFS LC-40	Dragon Spacecraft Qualification Unit	0	LEO	SpaceX	Success	Failure (parachute)
2010-12-08	15:43:00	F9 v1.0 B0004	CCAFS LC-40	Dragon demo flight C1, two CubeSats, barrel of Brouere cheese	0	LEO (ISS)	NASA (COTS) NRO	Success	Failure (parachute)
2012-05-22	07:44:00	F9 v1.0 B0005	CCAFS LC-40	Dragon demo flight C2	525	LEO (ISS)	NASA (COTS)	Success	No attempt
2012-10-08	00:35:00	F9 v1.0 B0006	CCAFS LC-40	SpaceX CRS-1	500	LEO (ISS)	NASA (CRS)	Success	No attempt
2013-03-01	15:10:00	F9 v1.0 B0007	CCAFS LC-40	SpaceX CRS-2	677	LEO (ISS)	NASA (CRS)	Success	No attempt

Total Payload Mass

sum_payload: % %sql

45596

```
SELECT SUM (payload_mass__kg_) AS
SUM_PAYLOAD FROM SPACEXDATASET
WHERE customer = 'NASA (CRS)';
```

sum_payload

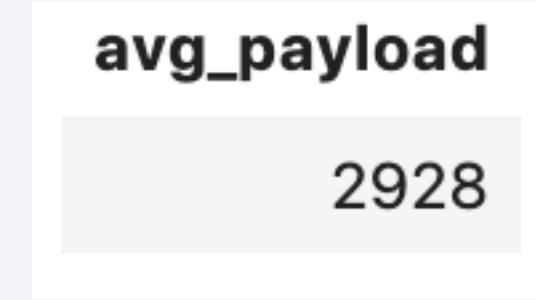
45596

Average Payload Mass by F9 v1.1

avg_payload: % %sql

2928

```
SELECT AVG (payload_mass_kg_) AS
AVG_PAYLOAD FROM SPACEXDATASET
WHERE BOOSTER_VERSION ='F9 v1.1';
```



First Successful Ground Landing Date

min_date: % %sql

2015-12-22 `SELECT MIN(DATE) AS MIN_DATE FROM
SPACEXDATASET WHERE landing_outcome
= 'Success (ground pad)';`

min_date

2015-12-22

Successful Drone Ship Landing with Payload between 4000 and 6000

```
% %sql
```

```
SELECT booster_version FROM SPACEXDATASET WHERE
payload_mass_kg_ > '4000' AND payload_mass_kg_ <
'6000' AND landing_outcome = 'Success (drone ship)';
```

booster_version

F9 FT B1022

F9 FT B1026

F9 FT B1021.2

F9 FT B1031.2

Total Number of Successful and Failure Mission Outcomes

Success: 100

% %sql

```
SELECT COUNT(*) AS SUCCESS FROM  
SPACEXDATASET WHERE  
mission_outcome LIKE 'Success%';
```

success

100

Boosters Carried Maximum Payload

- This query returns the booster versions that carried the highest payload mass of 15600 kg.
- These booster versions are very similar.
- This likely indicates payload mass correlates with the booster version that is used.

% %sql

```
SELECT booster_version,(SELECT  
MAX(payload_mass_kg_) FROM  
SPACEXDATASET) AS MAX_Booster  
FROM SPACEXDATASET ;
```

booster_version	max_booster
F9 v1.0 B0003	15600
F9 v1.0 B0004	15600
F9 v1.0 B0005	15600
F9 v1.0 B0006	15600
F9 v1.0 B0007	15600
F9 v1.1 B1003	15600
F9 v1.1	15600
F9 v1.1 B1011	15600
F9 v1.1 B1010	15600
F9 v1.1 B1012	15600
F9 v1.1 B1013	15600
F9 v1.1 B1014	15600

2015 Launch Records

% %sql

```
SELECT Date, booster_version, launch_site, landing__outcome FROM
SPACEXDATASET WHERE landing__outcome = 'Failure (drone ship)' AND YEAR(Date)
= 2015;
```

DATE	booster_version	launch_site	landing__outcome
2015-01-10	F9 v1.1 B1012	CCAFS LC-40	Failure (drone ship)
2015-04-14	F9 v1.1 B1015	CCAFS LC-40	Failure (drone ship)

Rank Landing Outcomes Between 2010-06-04 and 2017-03-20

% %sql

```
SELECT landing_outcome FROM SPACEXDATASET WHERE
Date > '2010-06-04' AND Date < '2017-03-20' GROUP BY
landing_outcome ORDER BY COUNT(landing_outcome)
DESC;
```

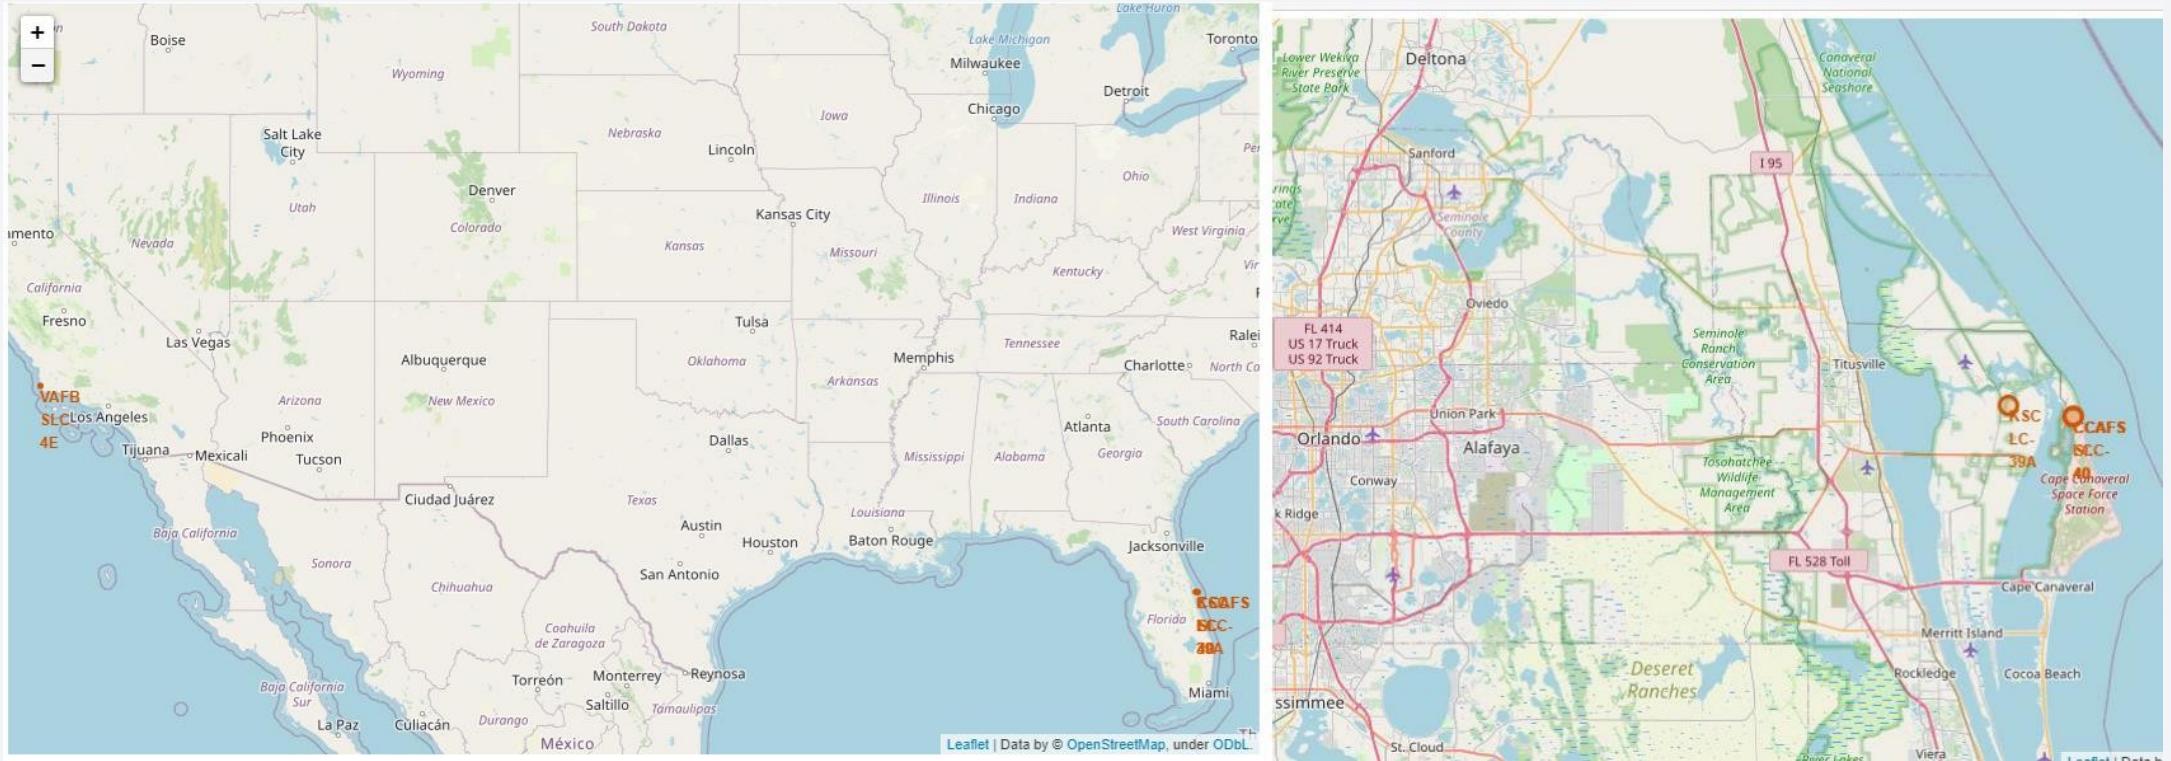
landing_outcome
No attempt
Failure (drone ship)
Success (drone ship)
Controlled (ocean)
Success (ground pad)
Uncontrolled (ocean)
Failure (parachute)
Precluded (drone ship)

A nighttime satellite view of Earth from space, showing city lights and auroras.

Section 3

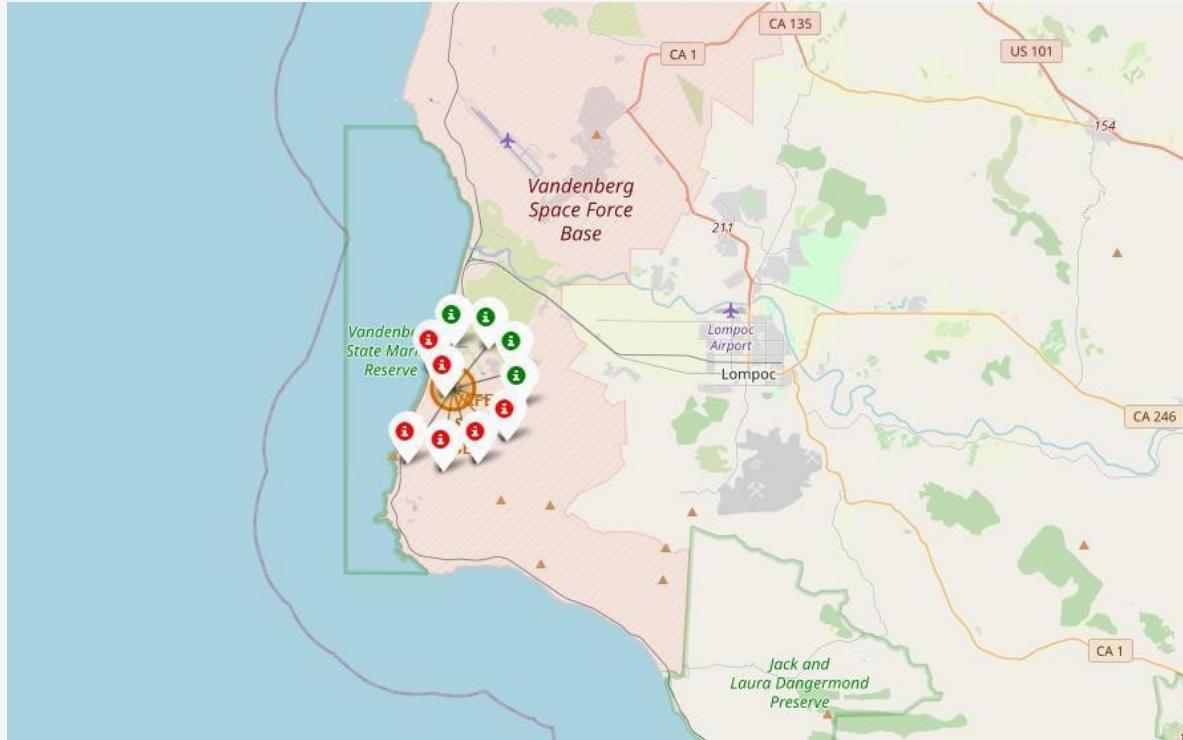
Launch Sites Proximities Analysis

Launch Site Locations



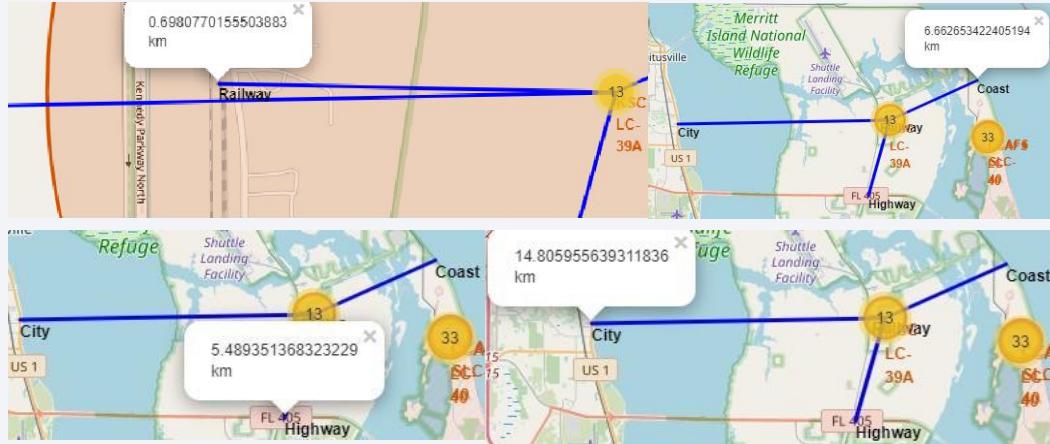
The left map shows all launch sites relative US map. The right map shows the two Florida launch sites since they are very close to each other. All launch sites are near the ocean.

Color Coded Launch Markers



Clusters on Folium map can be clicked on to display each successful landing (green icon) and failed landing (red icon). In this example VAFB SLC-4E shows 4 successful landings and 6 failed landings.

Key Location Proximities



Using KSC LC-39A as an example, launch sites are very close to railways for large part and supply transportation. Launch sites are close to highways for human and supply transport. Launch sites are also close to coasts and relatively far from cities so that launch failures can land in the sea to avoid rockets falling on densely populated areas.

Section 5

Predictive Analysis (Classification)

Classification Accuracy

All models had virtually the same accuracy on the test set at 83.33% accuracy. It should be noted that test size is small at only sample size of 18.

This can cause large variance in accuracy results, such as those in Decision Tree Classifier model in repeated runs. We likely need more data to determine the best model.

Confusion Matrix



Correct predictions are on a diagonal from top left to bottom right.

Since all models performed the same for the test set, the confusion matrix is the same across all models. The models predicted 12 successful landings when the true label was successful landing.

The models predicted 3 unsuccessful landings when the true label was unsuccessful landing.

The models predicted 3 successful landings when the true label was unsuccessful landings (false positives). Our models over predict successful landings.

Conclusions

- Our task: to develop a machine learning model for Space Y who wants to bid against SpaceX
- The goal of model is to predict when Stage 1 will successfully land to save ~\$100 million USD
- Used data from a public SpaceX API and web scraping SpaceX Wikipedia page
- Created data labels and stored data into a DB2 SQL database
- Created a dashboard for visualization
- We created a machine learning model with an accuracy of 83%
- Allon Mask of SpaceY can use this model to predict with relatively high accuracy whether a launch will have a successful Stage 1 landing before launch to determine whether the launch should be made or not
- If possible more data should be collected to better determine the best machine learning model and improve accuracy

Appendix

GitHub Repository URL:

<https://github.com/taeefnajib/IBM-Applied-Data-Science-Capstone>

Instructors:

Rav Ahuja, Alex Aklson, Aije Egwaikhide, Svetlana Levitan, Romeo Kienzler, Polong Lin, Joseph Santarcangelo, Azim Hirjani, Hima Vasudevan, Saishruthi Swaminathan, Saeed Aghabozorgi, Yan Luo

Special Thanks to All Instructors:

<https://www.coursera.org/professional-certificates/ibm-data-science?#instructors>

Thank you!

